

CS434 – Data Base Theory and Design

Project #5

Team Database Application (TDA): Part 5 – Querying and Schema Tuning

Team

Lipika Baniya | 800794205 | lbaniya@siue.edu

The domain I would like to manage with the TDA is **Washington DC Crime Datasets 2024** by the District of Columbia Metropolitan Police Department (MPD).

General Nature of application

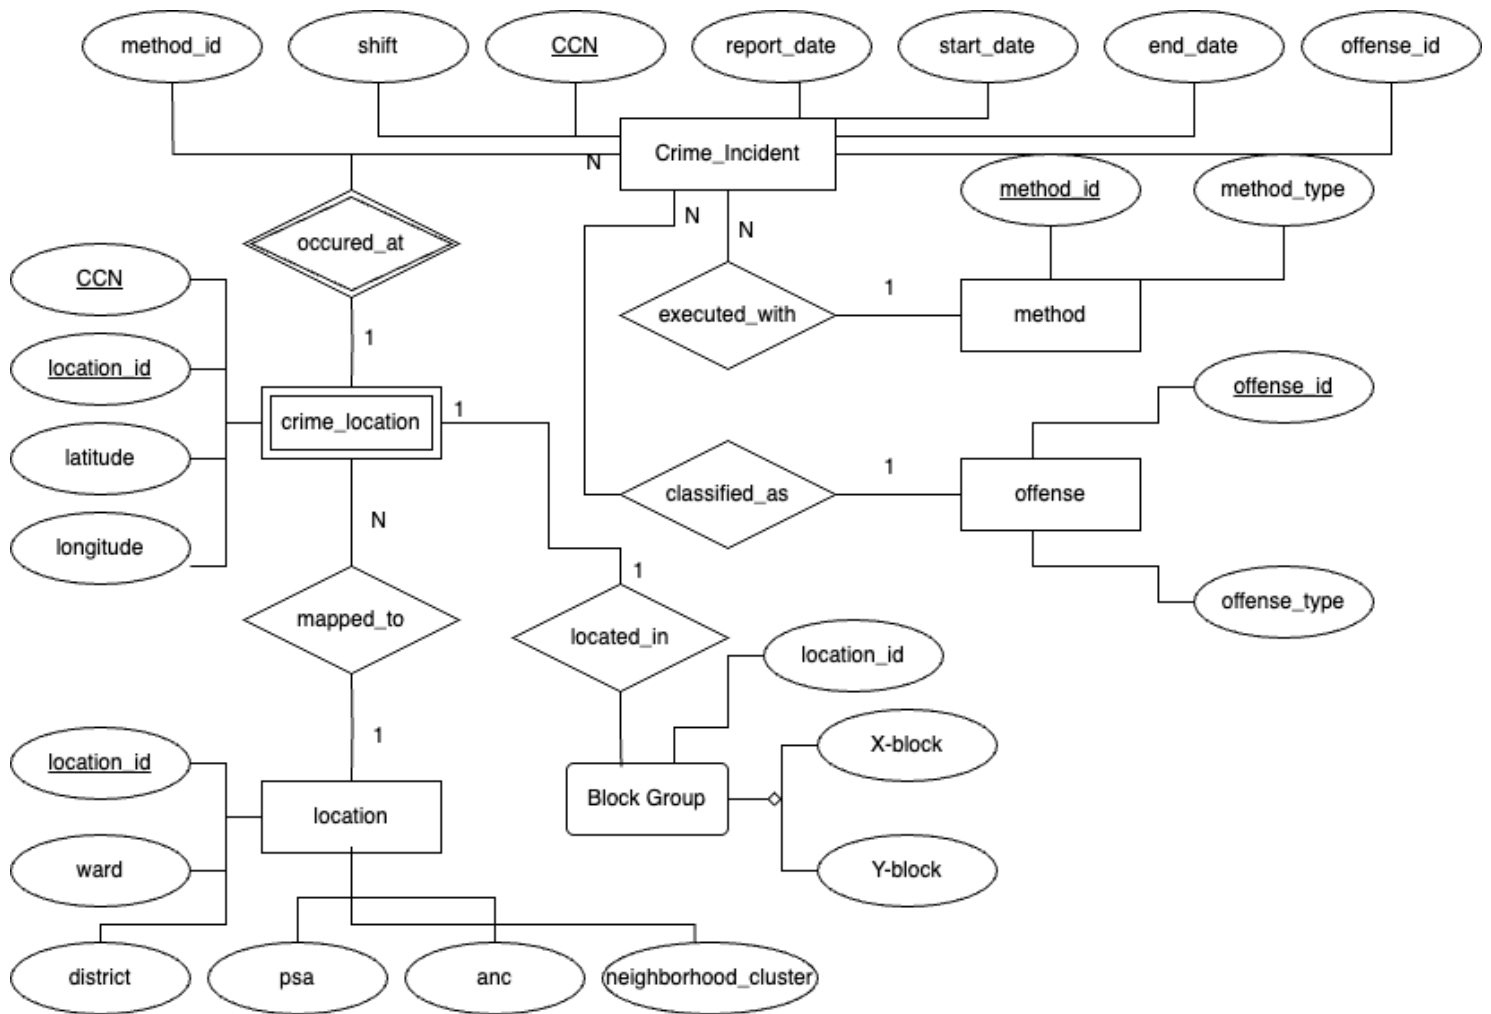
The main goal of an Entity Relationship Diagram (ER Diagram) is to explain the relationship between entities; it is a structural design of the database. Through the help of specialized symbols, it helps to define the relationship between entities. It is based on three main principles entities, attributes and relationships, these help to design the database that would be required before implementing the database. It is a systematic process to design a database as it would require analyzing all requirements.

About Data

Washington, D.C. has been facing significant challenges in ensuring public safety due to the varying and growing crime rates in different neighborhoods and time periods. It is important for law enforcement agencies to understand when and where crimes occur so that it can respond efficiently and allocate limited resources wisely. Imagine a robust database system that is designed to handle this task effectively, because without a data-driven approach and structured database, policing efforts may remain reactive, which would result in delays or gaps in coverage in high-risk areas. This database includes various entities, each representing a key component of crime data management.

In the previous project, I created the schema and inserted data into PostgreSQL based on the ER diagram. I used PostgreSQL to CREATE TABLE command and inserted data values and types for each entity's attributes. I have attached screenshots of the entire dataset and separate files for each table relation which were included in Project 4.

ER Diagram



1. Queries

1.1. Neighborhood Clusters with highest number of reported theft incidents

Question: Which neighborhood clusters have the highest number of reported theft incidents?

Query
Query History

```

1 SELECT l.neighborhood_cluster, COUNT(*) AS theft_incidents
2 FROM crime_incident ci
3 JOIN offense o ON ci.offense_id = o.offense_id
4 JOIN crime_location cl ON cl.ccn = ci.ccn
5 JOIN location l ON l.location_id = cl.location_id
6 WHERE o.offense_type ILIKE '%theft%'
7 GROUP BY l.neighborhood_cluster
8 ORDER BY theft_incidents DESC
    
```

Data Output
Messages
Notifications

SQL

	neighborhood_cluster <small>character varying (100)</small>	theft_incidents <small>bigint</small>
1	Cluster 2	1781
2	Cluster 25	1635
3	Cluster 8	1555
4	Cluster 23	1446
5	Cluster 3	1324
6	Cluster 18	1209
7	Cluster 6	1196
8	Cluster 21	1021
9	Cluster 26	980
10	Cluster 7	917
11	Cluster 22	852
12	Cluster 17	787
13	Cluster 39	696
14	Cluster 1	646
15	Cluster 34	605
16	Cluster 27	595
17	Cluster 4	565
18	Cluster 33	529

Figure 1: Neighborhood Clusters with highest number of reported theft incidents

Figure 1 shows the result of the query, which lists neighborhood clusters along with the number of theft-related incidents.

Explanation

To identify which neighborhood clusters experienced the most theft incidents, I used the following tables

- Crime_incident: this table is logbook of all crimes, including the offense_id
- Offense: provides the offense type (e.g. theft, assault)

- Crime_location links each incident to the neighborhood cluster
- Location: includes the detailed geographic attributes, including neighborhood cluster

The query:

- WHERE clause with LIKE '%theft%' filters only those crimes that are theft related.
- Neighborhood_cluster is grouped to count the number of thefts in each.
- Orders the result in descending order to show the clusters with the highest theft counts first.

1.2. Major time for crime to occur

Question: At what time during the day does the crime occur the most?

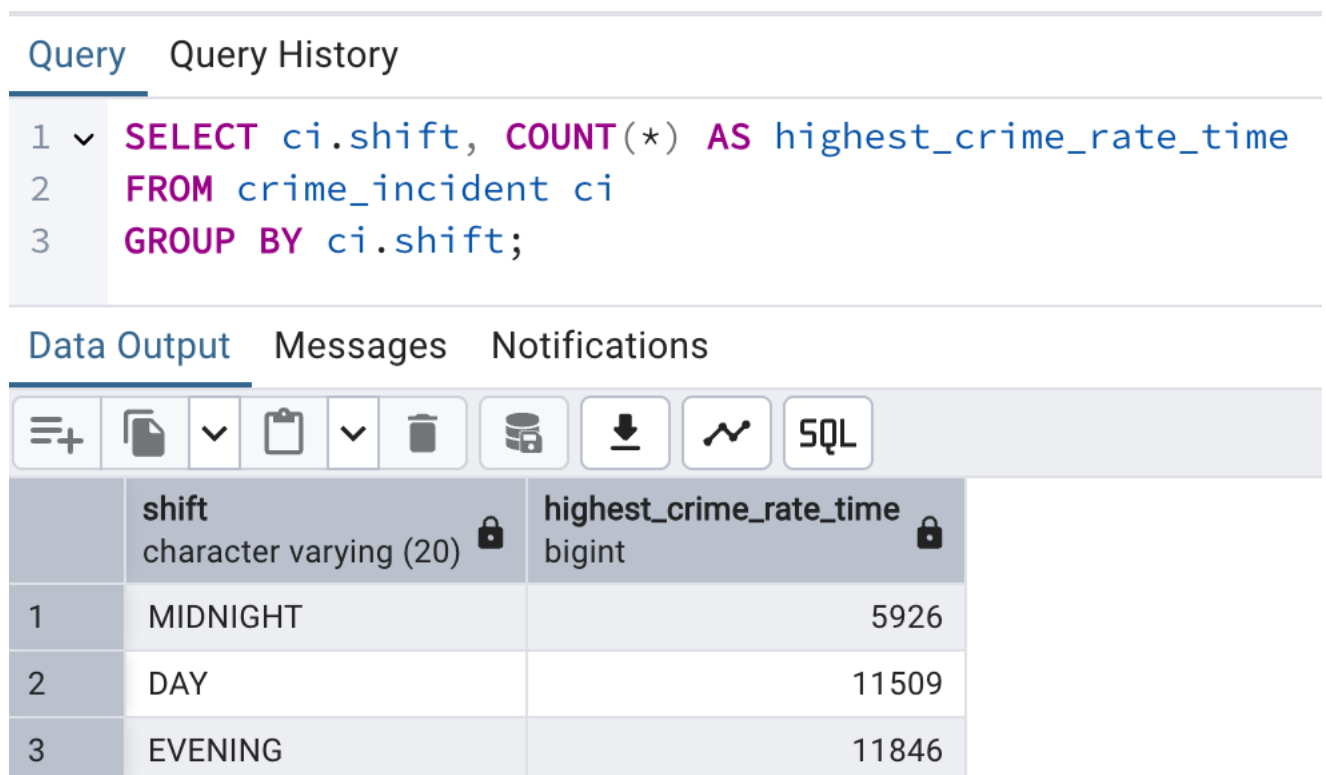


Figure 2: Major time for crime to occur

Figure 2 shows the result of the query, which displays the major time during the day does the crime occur the most.

Explanation

To find the major time for the crime to occur, I used the following tables

- Crime_incident: this table includes all crime details, including shift.

The query:

- The shift is grouped to show at what time the crime occurs the most.

1.3. Top Crime Type in Each Block Group

Question: What is the major crime in each block group?

Query

Query History

1

▼

WITH offense_counts AS (

2

SELECT

3

o.offense_type,

4

bl.x_block,

5

bl.y_block,

6

COUNT(*) AS offense_count,

7

ROW_NUMBER() OVER (

8

PARTITION BY bl.x_block, bl.y_block

9

ORDER BY COUNT(*) DESC

10

) AS rn

11

FROM crime_incident ci

12

JOIN offense o ON o.offense_id = ci.offense_id

13

JOIN crime_location cl ON cl.ccn = ci.ccn

14

JOIN block_group bl ON bl.location_id = cl.location_id

15

GROUP BY o.offense_type, bl.x_block, bl.y_block

16

)

17

SELECT offense_type, x_block, y_block, offense_count

18

FROM offense_counts

19

WHERE rn = 1;

20

Data Output

Messages

Notifications

≡+

📄

▼

📋

▼

🗑

🗄

⬇

📈

SQL

	offense_type character varying (100) 🔒	x_block numeric (10,2) 🔒	y_block numeric (10,2) 🔒	offense_count bigint 🔒
1	THEFT F/AUTO	390362.15	140808.01	66
2	THEFT F/AUTO	390448.11	140142.05	66
3	THEFT F/AUTO	390546.36	140878.18	66
4	THEFT F/AUTO	390591.47	140471.68	66
5	THEFT F/AUTO	390601.63	140260.37	66
6	THEFT F/AUTO	390602.16	140800.71	66
7	THEFT F/AUTO	390619.79	139840.73	66
8	THEFT F/AUTO	390689.92	139895.64	66
9	THEFT F/AUTO	390695.77	140453.52	66
10	THEFT F/AUTO	390815.95	139989.40	66

Figure 3: Top Crime Type in Each Block Group

Figure 3 displays the major crime in each block group.

Explanation

To find the major crime in each block group, the following tables were used:

- Crime_incident: the main record of incidents
- Offense: maps offense_id to a name like “THEFT F/AUTO”.
- Crime_location: links crimes to a location.
- Block_group: gives geographic blocks via x_block, y_block.

The query:

- WITH offense_counts AS is used to create Common Table Expression (CTE) to get temporary result named offense_counts.
- SELECT is used to find the Type of offense, X and Y coordinates of block_group and count of how many such offenses occurred.
- Row_number() for each (x_block, y_block) pair, ranks the most frequent offense in each block.
- Group_by is used to group offense_type and block_group to count number of times each offense happens in each block.
- The final selects query filters only the top-ranked offense type per block group.

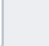









1.4.Incidents with “Gun” involved

Question: How many incidents involved a "Gun" as the method?

[Query](#) [Query History](#)

```
1  SELECT m.method_type, COUNT(*) AS total_gun_related_incidents
2  FROM crime_incident ci
3  JOIN method m ON m.method_id = ci.method_id
4  WHERE m.method_type ILIKE '%gun%'
5  GROUP BY m.method_type
6  ORDER BY total_gun_related_incidents DESC;
7
```

[Data Output](#) [Messages](#) [Notifications](#)



	method_type character varying (100)	total_gun_related_incidents bigint
1	GUN	2180

Figure 4: Incidents with "Gun" involved

Figure 4 shows the number of incidents that involved “gun” as the method.

Explanation

To find incidents that involved “Gun” as the method, I used the following the tables:

- Crime_incident: this table includes all crime details, including shift.
- Method: this table includes method details, including method type (e.g. “KNIFE”, “GUN”)

The query:

- JOIN is used to join crime_incident to the method table using method_id.
- WHERE ILIKE “%gun%” is used to find methods that use gun.
- GROUP BY is used to count the number of incidents that involved gun.
- ORDER BY is used to display in descending order.

1.5. Crime Reported Late

Question: How many crimes were reported more than 3 days after they occurred?

The screenshot shows a SQL query editor interface. At the top, there are tabs for 'Query' and 'Query History'. The 'Query' tab is active, displaying a SQL query with line numbers 1, 2, and 3. Below the query, there are tabs for 'Data Output', 'Messages', and 'Notifications'. The 'Data Output' tab is active, showing a table with one row and one column. The table has a header row with the column name 'late_reports' and data type 'bigint'. The first row of data shows the value '4006'.

```
1 SELECT COUNT(*) AS late_reports
2 FROM crime_incident ci
3 WHERE report_date - start_date > INTERVAL '3 days';
```

	late_reports bigint
1	4006

Figure 5: crimes that were reported late

Figure 5 shows the number of crimes reported after 3 days than they occurred.

Explanation

To find the number of reports that were reported late, I used the queries:

- Crime_incident: this table includes all crime details, including report_date, start_date.

The query:

- WHERE is used to subtract report_date – start_date to find interval of 3 days.

2.1. Insert

```
Query History
1  INSERT INTO offense (offense_id, offense_type)
2  VALUES (11, 'MURDER');
3
4  INSERT INTO method (method_id, method_type)
5  VALUES (11, 'Knife');
6
7  INSERT INTO location (location_id, ward, district, psa, ans, neighborhood_cluster)
8  VALUES (1798988, 2, '3D', 305, '2A', 'Columbia Heights');
9
10 INSERT INTO block_group (x_block, y_block, location_id)
11 VALUES (398765.12, 137654.33, 1);
12
13 INSERT INTO crime_incident (ccn, report_date, start_date, end_date, shift, offense_id, method_id)
14 VALUES (100000011, '2024-06-01', '2024-05-30', '2024-05-30', 'EVENING', 11, 11);
15
16 INSERT INTO crime_location (ccn, location_id, latitude, longitude)
17 VALUES (100000011, 1798988, 38.922, -77.032);
18
```

Data Output Messages Notifications

INSERT 0 1

Query returned successfully in 39 msec.

Figure 6: Insert Query

Query
Query History

```

1  SELECT * FROM crime_incident
2  WHERE ccn = '100000011';

```

Data Output
Messages
Notifications

Showing rows: 1 to 1
Page No: 1
of 1

	ccn [PK] character varying (20)	report_date timestamp without time zone	start_date timestamp without time zone	end_date timestamp without time zone	shift character varying (20)	offense_id integer	method_id integer
1	100000011	2024-06-01 00:00:00	2024-05-30 00:00:00	2024-05-30 00:00:00	EVENING	11	11

Figure 7: New tuple added to the DB

2.2.Update: Updating Shift

Query Query History

```
1  ✓ UPDATE crime_incident
2     SET shift = 'MORNING'
3     WHERE ccn = '100000011';
```

Data Output **Messages** Notifications

UPDATE 1

Query returned successfully in 88 msec.

Figure 8: Updating shift

Query

Query History

1

SELECT * FROM crime_incident

2

WHERE ccn = '100000011';

Data Output

Messages

Notifications

+

📄

▼

🗑️

▼

🗑️

📄

📄

SQL

Showing rows: 1 to 1

Page No: 1 of 1

	ccn [PK] character varying (20)	report_date timestamp without time zone	start_date timestamp without time zone	end_date timestamp without time zone	shift character varying (20)	offense_id integer	method_id integer
1	100000011	2024-06-01 00:00:00	2024-05-30 00:00:00	2024-05-30 00:00:00	MORNING	11	11

Figure 9: Updated record

2.3.Delete

Deleting the recently added dummy record from all the tables.

Query	Query History
1	-- Step 1: Delete from child table crime_location
2	✓ DELETE FROM crime_location
3	WHERE ccn = '100000011';
4	
5	-- Step 2: Delete from child table crime_incident
6	✓ DELETE FROM crime_incident
7	WHERE ccn = '100000011';
8	
9	-- Step 3: Delete from block_group
10	✓ DELETE FROM block_group
11	WHERE location_id = 1;
12	
13	-- Step 4: Delete from location
14	✓ DELETE FROM location
15	WHERE location_id = 1798988;
16	
17	-- Step 5: Delete from offense
18	✓ DELETE FROM offense
19	WHERE offense_id = 11 AND offense_type = 'MURDER';
20	
21	-- Step 6: Delete from method
22	✓ DELETE FROM method
23	WHERE method_id = 11 AND method_type = 'Knife';

Data Output	Messages	Notifications
-------------	----------	---------------

DELETE 1

Query returned successfully in 47 msec.

Figure 10: Delete the records

3. Create View

3.1.Creating the view table

A virtual table using existing schema has been created using create view as follows:

Query	Query History
<pre>1 CREATE VIEW Crime_view AS 2 SELECT ci.ccn, ci.report_date, o.offense_type AS Offense_Type, m.method_type AS method_description, l.neighborhood_cluster, 3 cl.latitude, cl.longitude, l.ward 4 FROM crime_incident as ci 5 JOIN offense o ON o.offense_id = ci.offense_id 6 JOIN method m ON m.method_id = ci.method_id 7 JOIN crime_location cl ON cl.ccn = ci.ccn 8 JOIN location l ON l.location_id = cl.location_id 9 JOIN block_group bg ON bg.location_id = cl.location_id;</pre>	
Data Output	Messages
CREATE VIEW	
Query returned successfully in 37 msec.	

Figure 11: Query for CREATE VIEW

3.2.The Crime View Table

Query

Query History

1

SELECT *

2

FROM Crime_view

3

LIMIT 20;

Data Output

Messages

Notifications

Showing rows: 1 to 20

Page No: 1 of 1

	ccn character varying (20)	report_date timestamp without time zone	offense_type character varying (100)	method_description character varying (100)	neighborhood_cluster character varying (100)	latitude numeric (9,6)	longitude numeric (9,6)	ward character varying (10)
1	24000457	2024-01-02 07:28:26	BURGLARY	OTHERS	Cluster 7	38.913711	-77.023968	2.0
2	24000745	2024-01-02 15:15:19	MOTOR VEHICLE THEFT	OTHERS	Cluster 7	38.908574	-77.023456	2.0
3	24001760	2024-01-04 14:57:40	THEFT F/AUTO	OTHERS	Cluster 7	38.912605	-77.022945	2.0
4	24001782	2024-01-04 15:38:41	MOTOR VEHICLE THEFT	OTHERS	Cluster 7	38.910385	-77.022942	2.0
5	24002167	2024-01-05 04:52:55	THEFT F/AUTO	OTHERS	Cluster 7	38.911123	-77.022433	2.0
6	24003708	2024-01-08 13:53:50	MOTOR VEHICLE THEFT	OTHERS	Cluster 7	38.912461	-77.021918	2.0
7	24004359	2024-01-09 20:08:07	THEFT F/AUTO	OTHERS	Cluster 7	38.908573	-77.026512	2.0
8	24004502	2024-01-10 01:44:20	THEFT/OTHER	OTHERS	Cluster 7	38.907238	-77.022432	2.0
9	24006974	2024-01-14 19:07:17	THEFT/OTHER	OTHERS	Cluster 7	38.909110	-77.023969	2.0
10	24007437	2024-01-15 18:02:23	THEFT/OTHER	OTHERS	Cluster 7	38.913348	-77.021915	2.0
11	24008377	2024-01-17 21:44:58	THEFT/OTHER	OTHERS	Cluster 7	38.910384	-77.021917	2.0
12	24008897	2024-01-18 21:26:04	THEFT/OTHER	OTHERS	Cluster 7	38.913348	-77.021915	2.0
13	24009631	2024-01-20 09:21:04	THEFT/OTHER	OTHERS	Cluster 7	38.912606	-77.023456	2.0
14	24010875	2024-01-23 01:33:19	THEFT/OTHER	OTHERS	Cluster 7	38.912605	-77.022431	2.0
15	24012325	2024-01-25 21:03:46	THEFT F/AUTO	OTHERS	Cluster 7	38.907909	-77.023966	2.0
16	24013283	2024-01-28 02:05:04	THEFT F/AUTO	OTHERS	Cluster 7	38.912605	-77.026512	2.0
17	24013333	2024-01-27 11:45:57	THEFT F/AUTO	OTHERS	Cluster 7	38.912606	-77.023456	2.0

Figure 12: Glimpse of the new Crime_view Table

3.3.Updating the View Table

Query	Query History
<pre>1 INSERT INTO Crime_view (ccn, report_date, offense_type, method_description, neighborhood_cluster, latitude, longitude, ward) 2 VALUES (1111111, '2025-06-30', 'ROBBERY', 'KNIFE', 'Cluster 8', 39.0, -56.90, 3);</pre>	
Data Output	Messages
ERROR: cannot insert into view "crime_view"	
Views that do not select from a single table or view are not automatically updatable.	
SQL state: 55000	
Detail: Views that do not select from a single table or view are not automatically updatable.	
Hint: To enable inserting into the view, provide an INSTEAD OF INSERT trigger or an unconditional ON INSERT DO INSTEAD rule.	

Figure 13: Updating View Table

The update view table query does not work and returned error; This is because the updated view cannot be inserted as it involves joins, and these are typically not updatable because the database cannot determine how to properly distribute new data across the base tables.

4. Indexing

4.1. Without indexing

For the below query, the query time was 163 msec.

[Query](#) [Query History](#)

1

2

3

4

5

6

7

8

9

```
SELECT l.neighborhood_cluster, COUNT(*) AS theft_count
FROM crime_incident ci
JOIN offense o ON ci.offense_id = o.offense_id
JOIN crime_location cl ON ci.CCN = cl.CCN
JOIN location l ON cl.location_id = l.location_id
WHERE o.offense_type ILIKE '%theft%' -- case-insensitive match
GROUP BY l.neighborhood_cluster
ORDER BY theft_count DESC;
```

Data Output [Messages](#) [Notifications](#)

Successfully run. Total query runtime: 163 msec.
46 rows affected.

4.2. Creating index

[Query](#) [Query History](#)

1

```
CREATE INDEX idx_crime_incident_ccn ON crime_incident (CCN);
```

Data Output [Messages](#) [Notifications](#)

CREATE INDEX

Query returned successfully in 186 msec.

[Query](#) [Query History](#)

1

2

```
CREATE INDEX idx_crime_location_ccn ON crime_location(ccn);
```

Data Output [Messages](#) [Notifications](#)

CREATE INDEX

Query returned successfully in 140 msec.

4.3. With indexing

With indexing, the query time dropped to 158 msec

Query	Query History
<pre>1 SELECT l.neighborhood_cluster, COUNT(*) AS theft_count 2 FROM crime_incident ci 3 JOIN offense o ON ci.offense_id = o.offense_id 4 JOIN crime_location cl ON ci.CCN = cl.CCN 5 JOIN location l ON cl.location_id = l.location_id 6 WHERE o.offense_type ILIKE '%theft%' -- case-insensitive match 7 GROUP BY l.neighborhood_cluster 8 ORDER BY theft_count DESC; 9</pre>	
Data Output	Messages
	<p>Successfully run. Total query runtime: 158 msec. 46 rows affected.</p>